

Introduction to Artificial Neural Networks

Hao Ji, Data Scientist

USC Center for Advanced Research Computing

Content

Introduction

Neural
Networks

Applications

PyTorch

Section 1: Introduction to Deep Learning

Section 2: Neural Networks

Section 3: Applications

Section 4: Building Neural Networks with PyTorch

Introduction

Neural
Networks

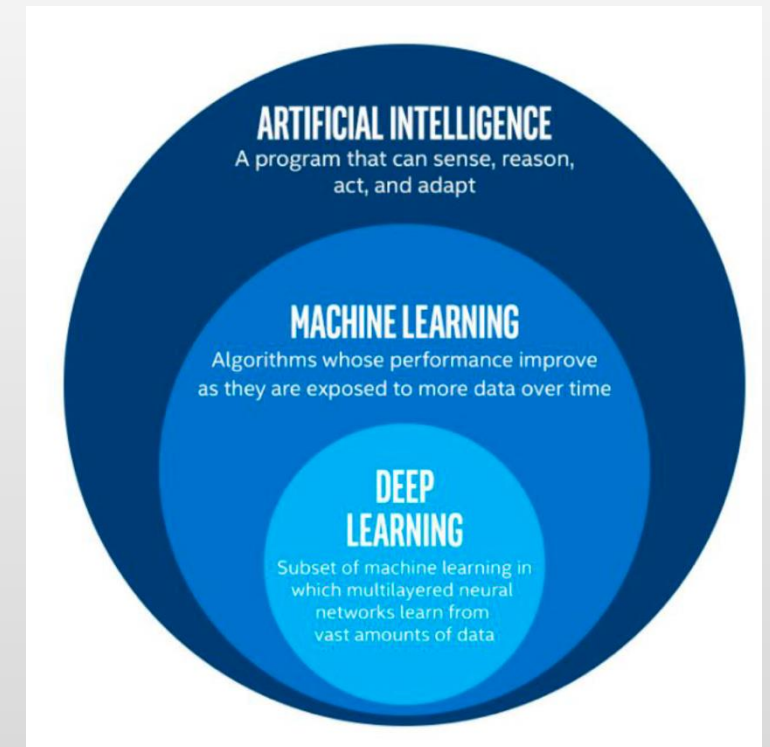
Applications

PyTorch

Introduction to Deep Learning

Deep Learning: subfield of traditional machine learning

- Inspired by the structure and function of the brain:
Artificial Neural Networks
- Computer vision: Tesla recognizing items on a street
- Text Generation: An algorithm trained to create a new Shakespeare piece
- Speech recognition
- Computer Games: AlphaGo

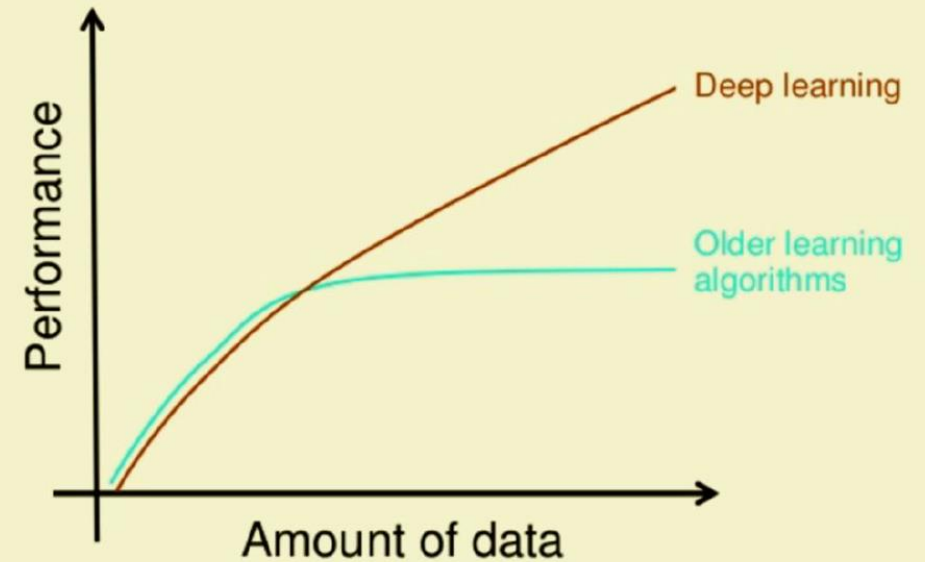


Introduction to Deep Learning

What drives the recent development of Deep Learning?

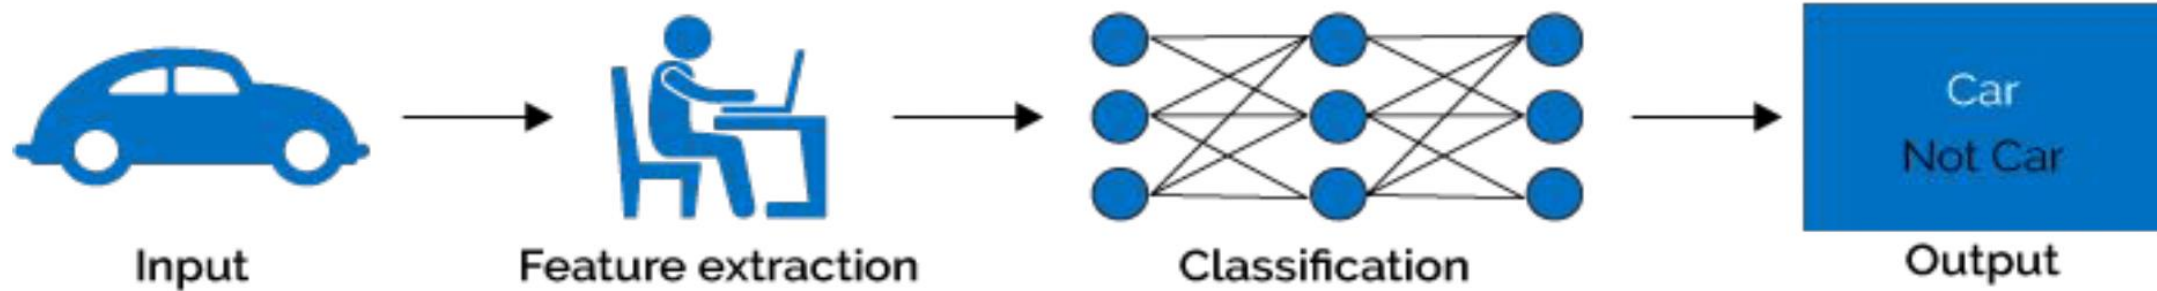
- Larger amounts of data available
- Data Storage
- Strong computation units such as GPU's

Why deep learning

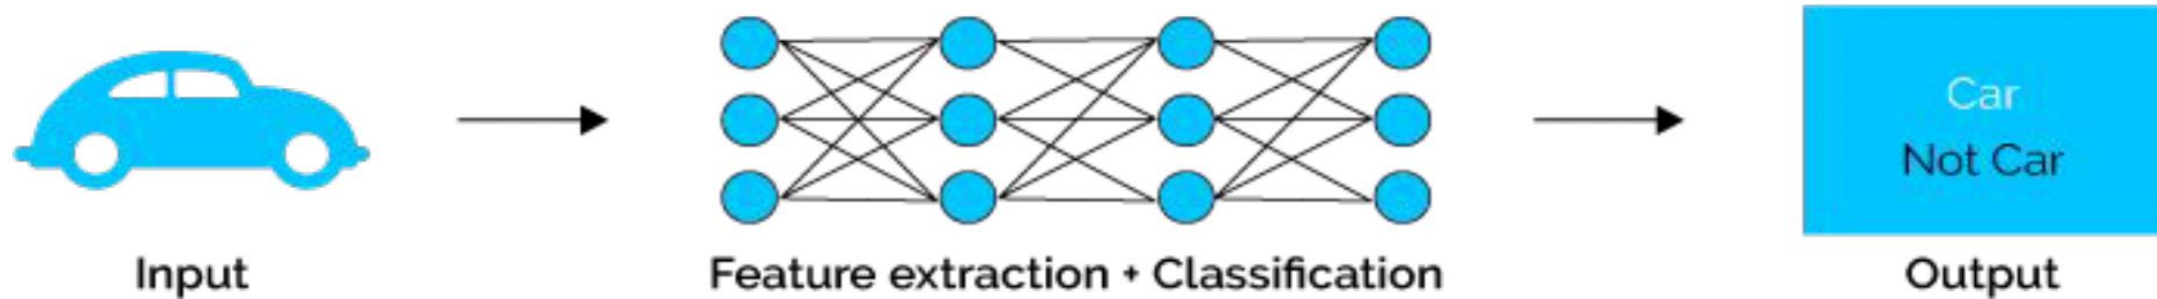


Introduction to Deep Learning

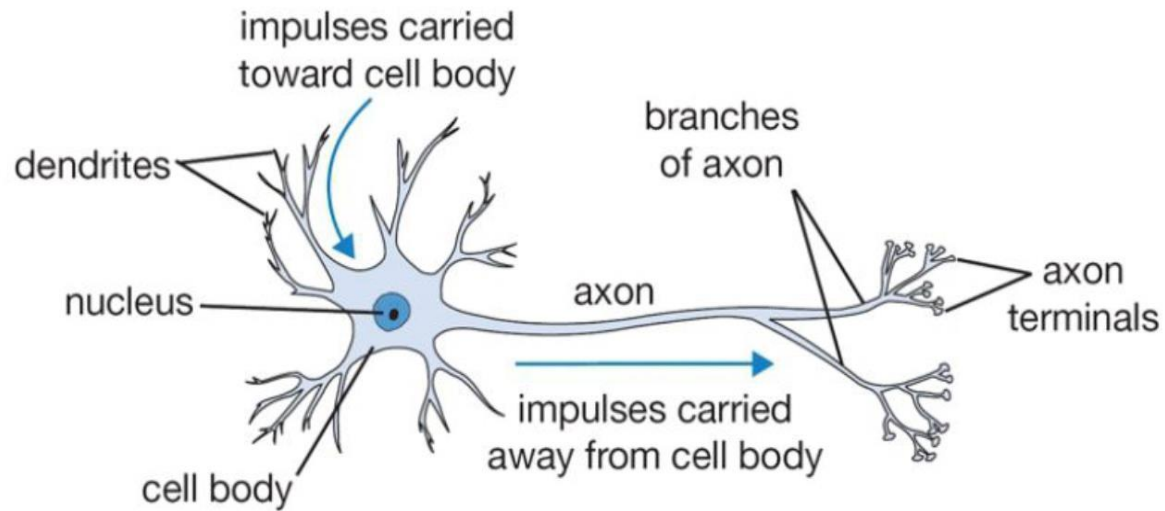
Machine Learning



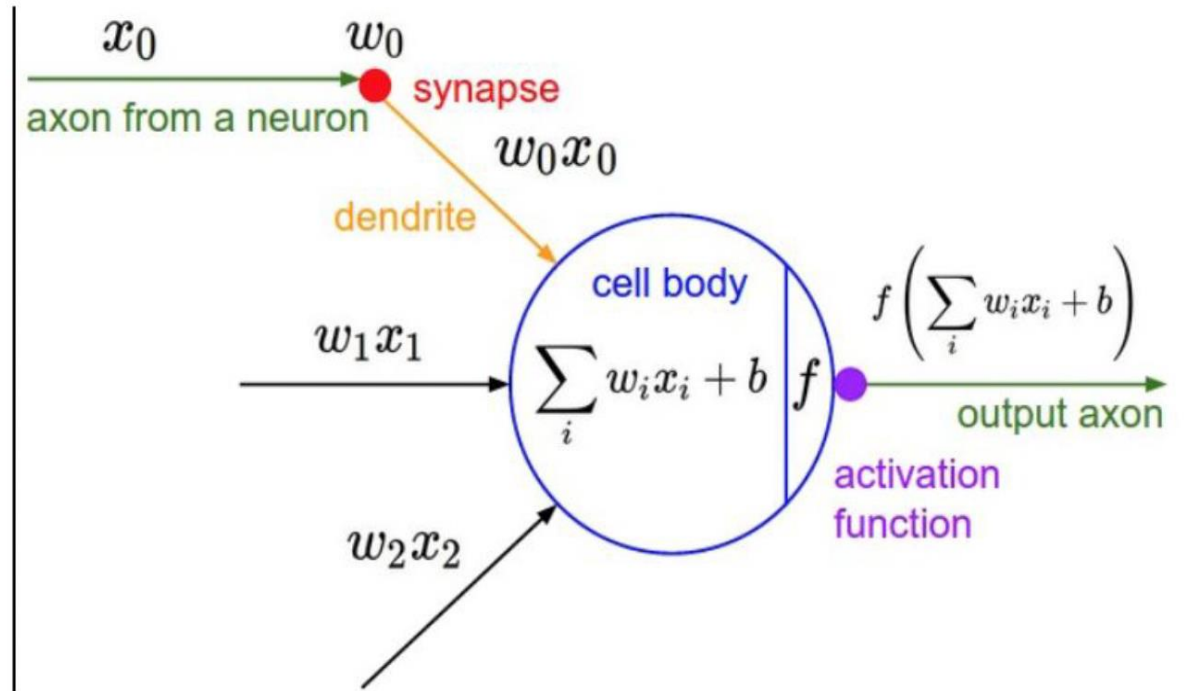
Deep Learning



Neural Networks



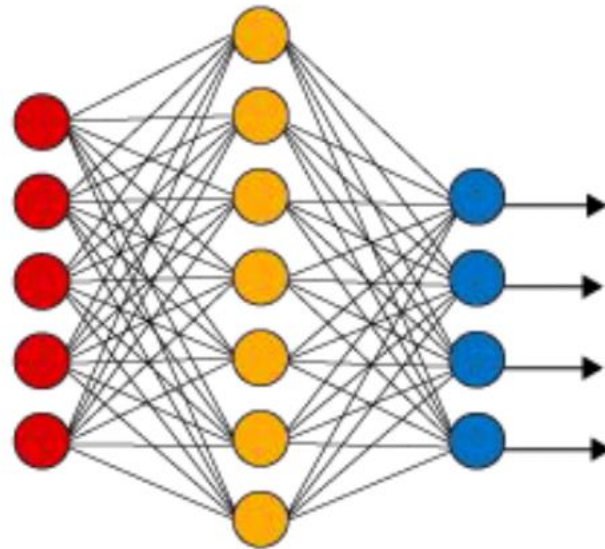
biological neuron



artificial neural networks

Neural Networks

Simple Neural Network

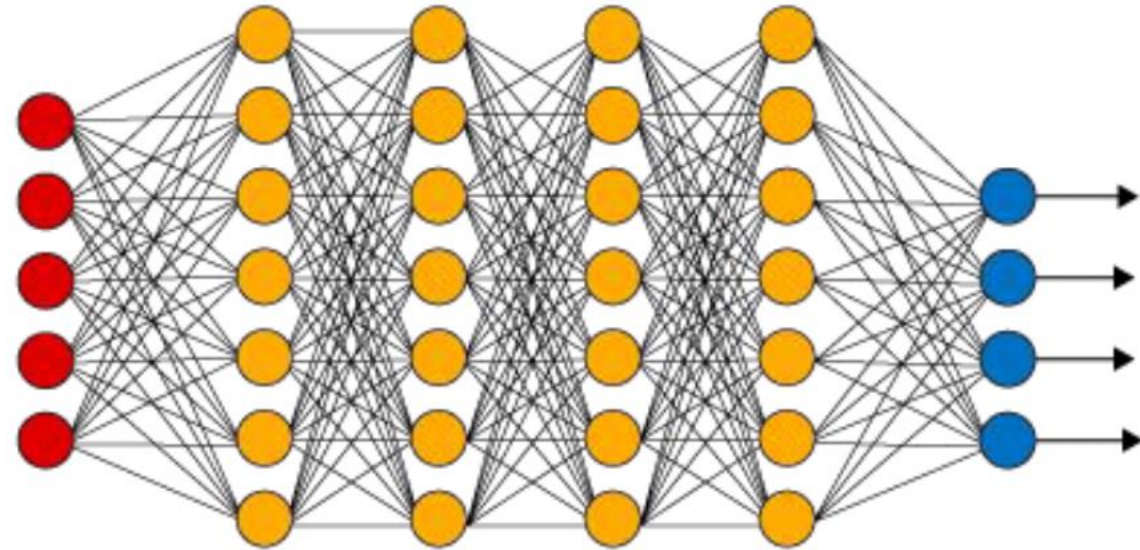


● Input Layer

● Hidden Layer

● Output Layer

Deep Learning Neural Network



A neural network (NN) has 3 types of layers:
Input layer Hidden layer Output layer

Deep neural networks (DNN) usually has more hidden layers
Still has same 3 types of layers

Building Neural Networks

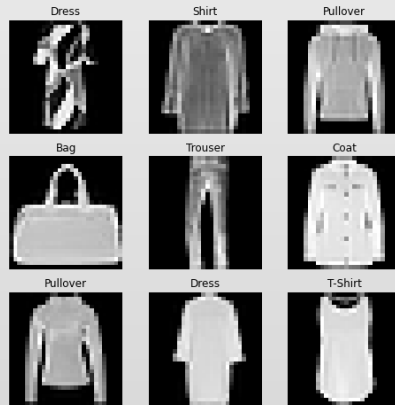
Task: Predict if an input image belongs to one of the following classes: T-shirt/top, Trouser, Pullover, Dress, Coat, Sandal, Shirt, Sneaker, Bag, or Ankle boot.



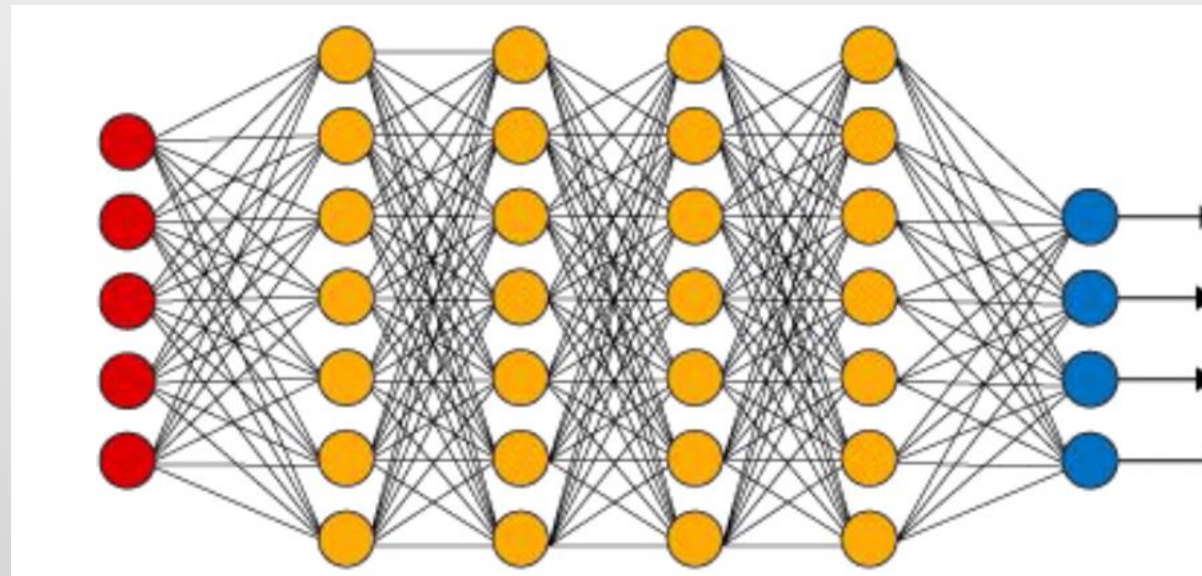
FasionMNIST Dataset

Fashion-MNIST is a dataset comprising of 28×28 grayscale images of 70,000 fashion products from 10 categories, with 7,000 images per category. The training set has 60,000 images and the test set has 10,000 images.

Building Neural Networks




Input X: (28 * 28)





Make Predictions based on Logits

Building Neural Networks



$28 * 28 = 784$





$$Z_1^{[1]} = w_{11}^{[1]}x_1 + w_{12}^{[1]}x_2 + w_{13}^{[1]}x_3 + b_1^{[1]}$$

$$= 0.5 * 0.4 + 0.1 * 0.2 + 0.8 * 0.1 + 0.2 = 0.5$$

$$a_1^{[1]} = f(0.5) = \text{ReLU}(0.5) = 0.5$$

Weights are initialized randomly

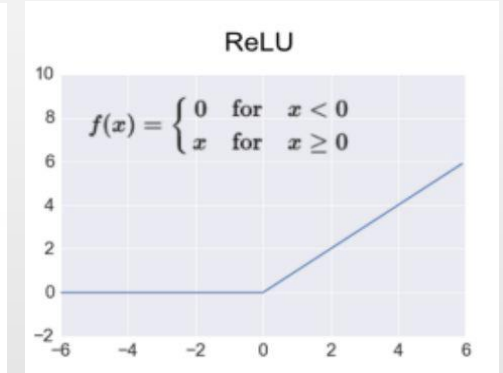
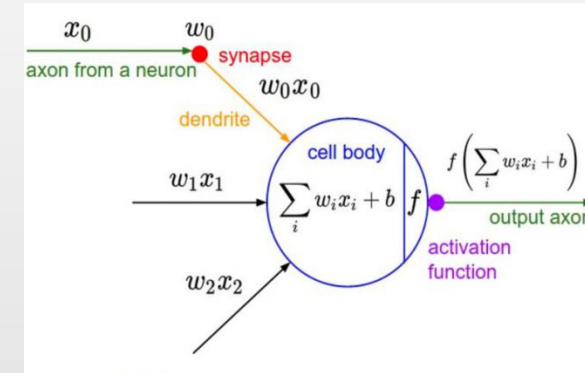
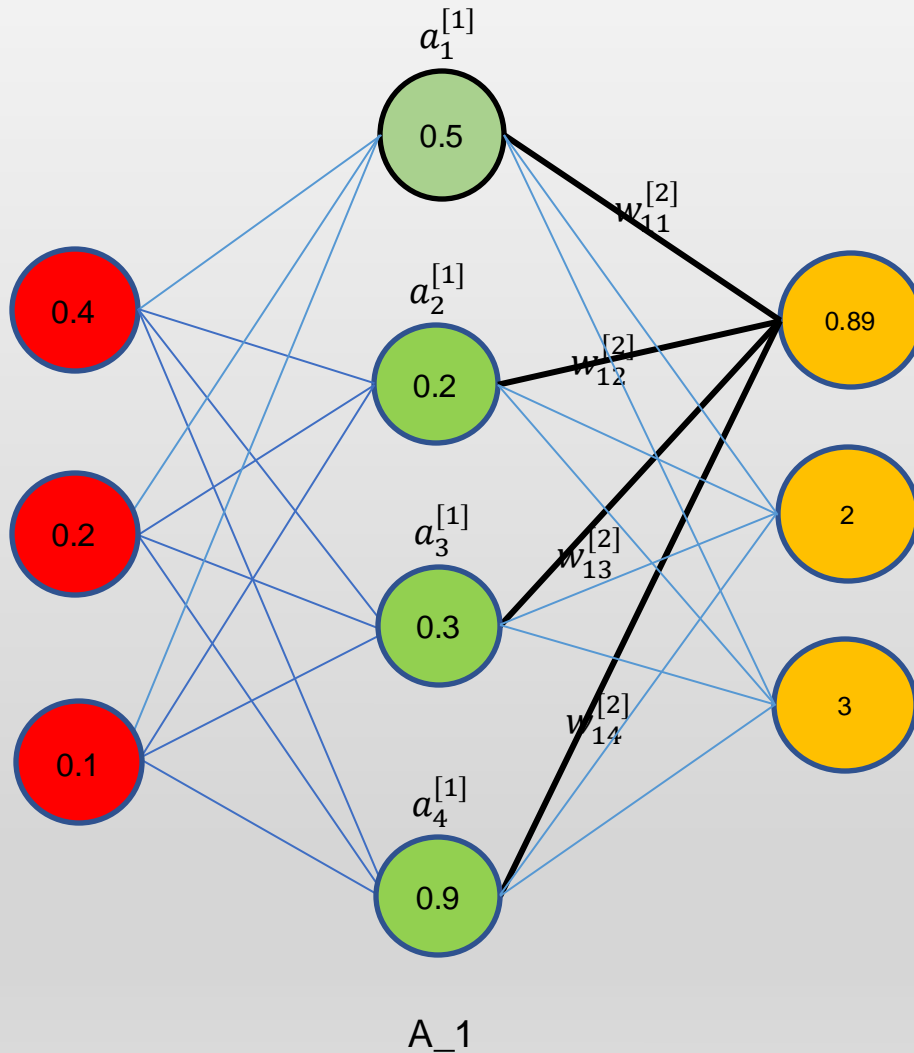
Building Neural Networks



$$\begin{aligned} Z_1^{[1]} &= w_{11}^{[1]} x_1 + w_{12}^{[1]} x_2 + w_{13}^{[1]} x_3 + b_1^{[1]} \\ &= 0.5 * 0.4 + 0.1 * 0.2 + 0.8 * 0.1 + 0.2 = 0.5 \end{aligned}$$

$$a_1^{[1]} = f(0.5) = \text{ReLU}(0.5) = 0.5$$

Building Neural Networks



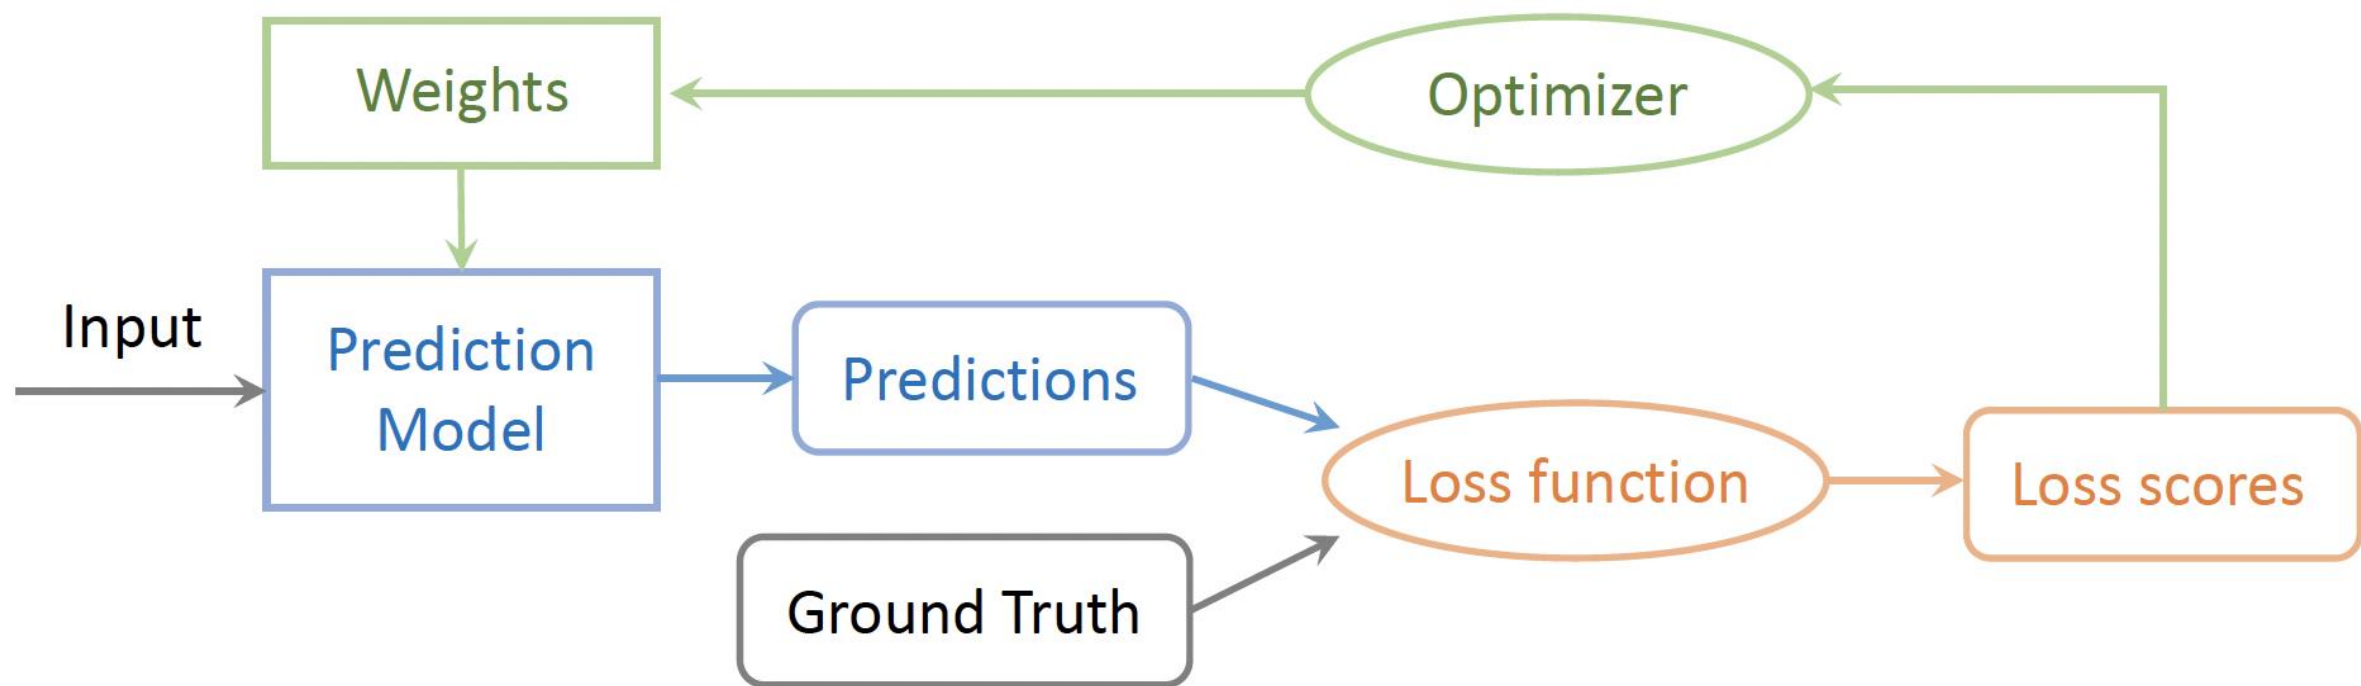
$$\begin{aligned}
 Z_1^{[2]} &= w_{11}^{[2]} a_1^{[1]} + w_{12}^{[2]} a_2^{[1]} + w_{13}^{[2]} a_3^{[1]} + w_{14}^{[2]} a_4^{[1]} + b_1^{[2]} \\
 &= 0.3 * 0.5 + 0.1 * 0.2 + 0.2 * 0.3 + 0.4 * 0.9 + 0.3 \\
 &= 0.89 \\
 a_1^{[2]} &= f(Z_1^{[2]}) = f(0.89) = \text{ReLU}(0.89) = 0.89
 \end{aligned}$$

Neural Networks

Three steps to training a neural network

- 1 **Forward propagation**: push example through the network to get a predicted output
- 2 **Compute the cost**: calculate the difference between predicted output and actual data
- 3 **Backward propagation**: push back the derivative of the error and apply to each weight, such that next time it will result in a lower error

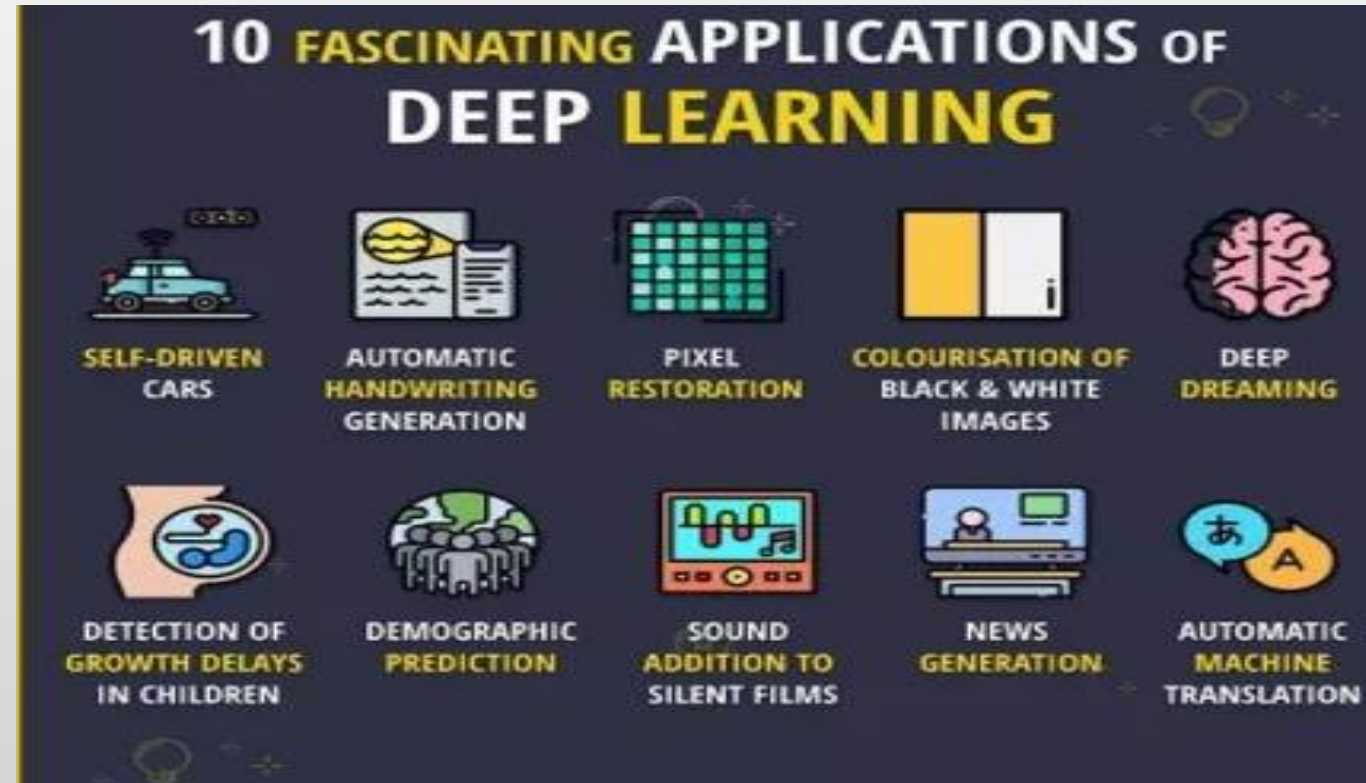
Training Pipeline



The training pipeline consists of choosing the prediction model, the loss function and the optimizer.

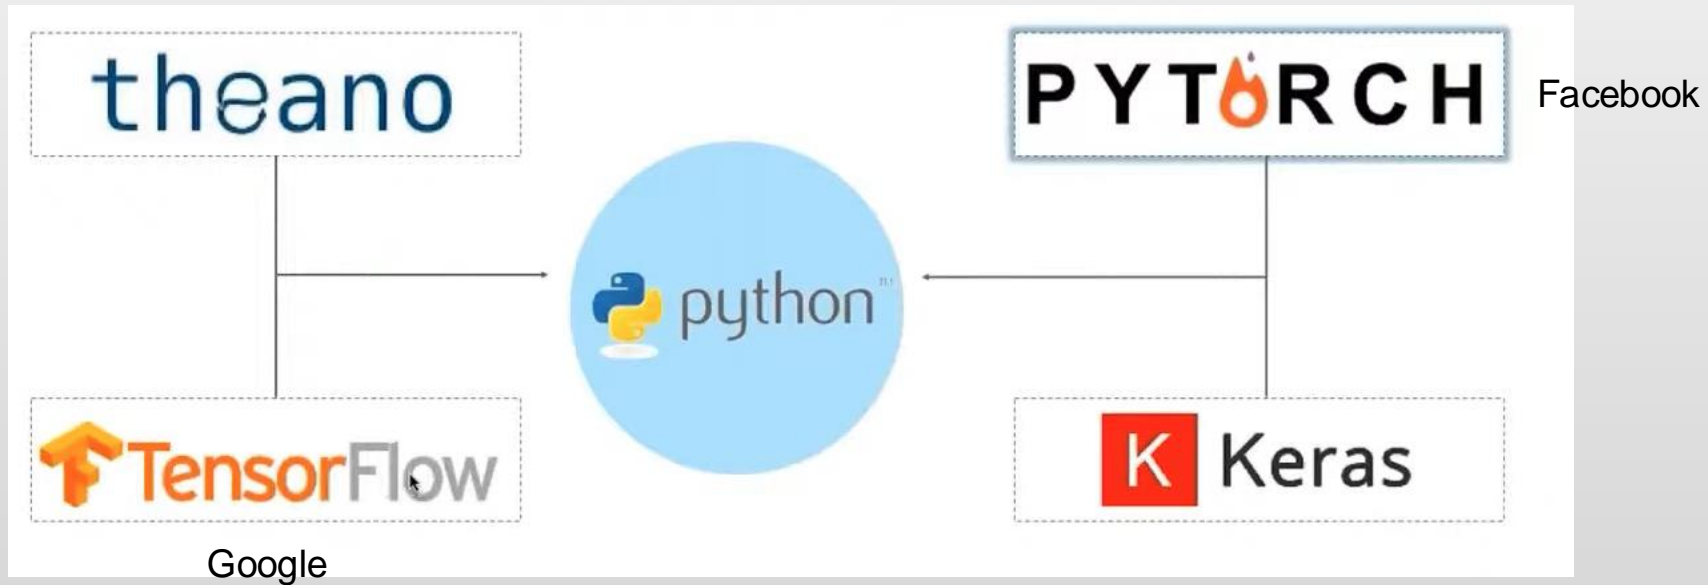
Once these choices are made, we can feed the input data and labels to start the training process.

Deep Learning Applications



[Deep Learning Applications](#)

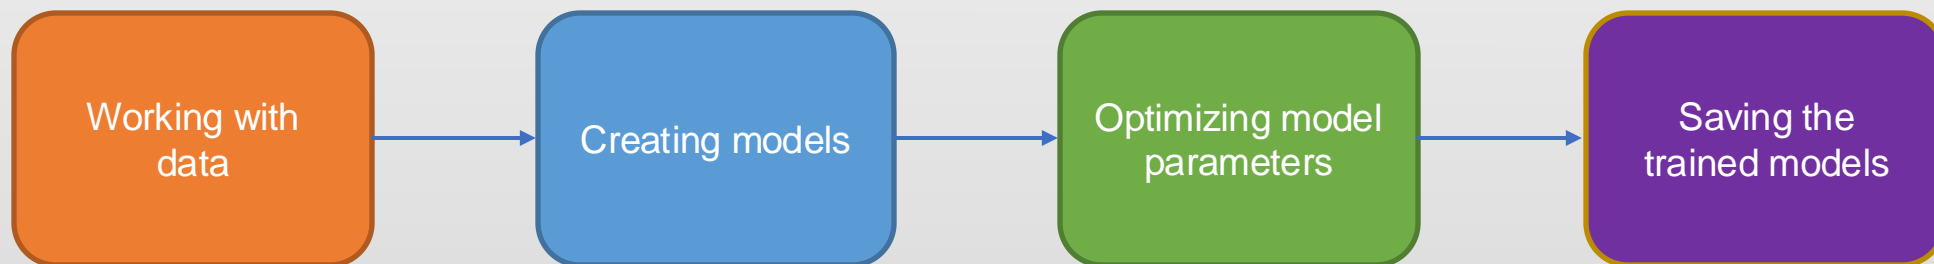
Neural Networks Packages



Pytorch Tensors

- Tensors are a specialized data structure similar to arrays and matrices
- PyTorch uses tensors to encode the inputs and outputs of a model, as well as model's parameters
- Can run on GPUs or other hardware accelerators
- Optimized for automatic differentiation

Machine Learning Workflows



Pytorch Tutorial

Predict if an input image belongs to one of the following classes: T-shirt/top, Trouser, Pullover, Dress, Coat, Sandal, Shirt, Sneaker, Bag, or Ankle boot.

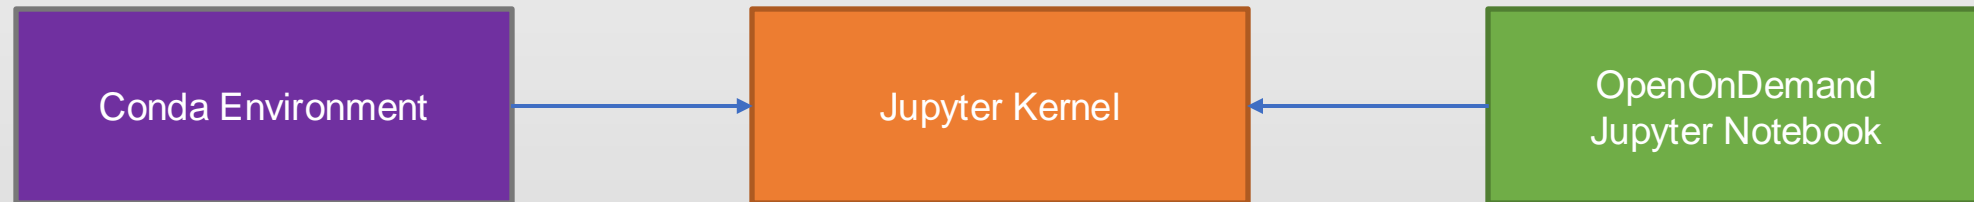


FasionMNIST Dataset

Using Anaconda & Jupyter Kernel

Let's build our Conda Environment First:

<https://github.com/uschpc/Building-NeuralNetworks.git>



CARC OnDemand

<https://www.carc.usc.edu>

Introduction

Neural
Networks

Applications

PyTorch

Search

Q

Latest News

► Services

► User Groups

▼ User Guides

► Quick Start Guides

► Project and Allocation Management

► HPC Systems

▼ CARC OnDemand

CARC OnDemand Overview

Data Management

Running Jobs with CARC OnDemand

Shell Access

Interactive Apps

► Research Data Management


► Advanced HPC Programming

► Life Sciences Computing

► Data Science

► Research Applications

► Cloud Computing

Center for Advanced Research Computing
Enabling scientific breakthroughs at scale

X D Y R

CARC OnDemand Overview

Last updated July 12, 2023

CARC's OnDemand service provides users graphical, browser-based access to the Discovery and Endeavour HPC clusters and their /home1, /project, /cryoem2, and /scratch1 directories. OnDemand offers:

- Easy file management
- Command line shell access
- Slurm job submission and management
- Interactive applications, including Jupyter notebooks and RStudio Server

OnDemand is available to all CARC users. To access OnDemand, you must belong to an active project in the [CARC user portal](#).

CARC OnDemand will only be accessible via a connection to either USC's Secure network or a USC VPN. Instructions for setting up a VPN connection can be found at the following links:

- [Windows](#)
- [MacOS](#)
- [Linux](#)

[Log in to CARC OnDemand](#)

We recommend using OnDemand in a private browser to avoid potential permissions issues related to your browser's cache. If you're using a private browser and still encounter permissions issues, please [submit a help ticket](#).

Table of Contents

- [Use cases](#)
- [Additional resources](#)

CARC OnDemand

Files

Jobs

Clusters


Interactive Apps

My Interactive Sessions

Help

Logged in as haoji

Log Out

Advanced Research Computing
Enabling scientific breakthroughs at scale

OnDemand provides an integrated, single access point for all of your HPC resources.

powered by
OPEN OnDemand

OnDemand version: v1.8.18

Pytorch Tutorial: Working with Data

PyTorch



Dataset: stores the samples and their corresponding labels TorchText, TorchVision, and TorchAudio

torchvision.datasets: CIFAR, COCO, FashionMNIST

DataLoader: wraps an iterable around the Dataset

Working with Data

```
import torch
from torch import nn
from torch.utils.data import DataLoader
from torchvision import datasets
from torchvision.transforms import ToTensor
```

Importing Modules

```
# Download training data from open datasets.
training_data = datasets.FashionMNIST(
    root="data",
    train=True,
    download=True,
    transform=ToTensor(),
)

# Download test data from open datasets.
test_data = datasets.FashionMNIST(
    root="data",
    train=False,
    download=True,
    transform=ToTensor(),
)
```

Define Training and Test Dataset

Working with Data

```
import torch
from torch import nn
from torch.utils.data import DataLoader
from torchvision import datasets
from torchvision.transforms import ToTensor
```

Importing Modules

```
# Download training data from open datasets.
training_data = datasets.FashionMNIST(
    root="data",
    train=True,
    download=True,
    transform=ToTensor(),
)

# Download test data from open datasets.
test_data = datasets.FashionMNIST(
    root="data",
    train=False,
    download=True,
    transform=ToTensor(),
)
```

Define Training and Test Dataset

Working with Data

Pass the Dataset as an argument to DataLoader

Wraps an iterable over dataset and supports automatic batching, sampling, shuffling and multiprocess data loading

```
batch_size = 64

# Create data loaders.
train_dataloader = DataLoader(training_data, batch_size=batch_size)
test_dataloader = DataLoader(test_data, batch_size=batch_size)

for X, y in test_dataloader:
    print(f"Shape of X [N, C, H, W]: {X.shape}")
    print(f"Shape of y: {y.shape} {y.dtype}")
    break
```

Define DataLoader

Creating Models

```
# Get cpu or gpu device for training.  
device = "cuda" if torch.cuda.is_available() else "cpu"  
print(f"Using {device} device")
```

Check if GPU is Available

```
# Define model  
class NeuralNetwork(nn.Module):  
    def __init__(self):  
        super(NeuralNetwork, self).__init__()  
        self.flatten = nn.Flatten()  
        self.linear_relu_stack = nn.Sequential(  
            nn.Linear(28*28, 512),  
            nn.ReLU(),  
            nn.Linear(512, 512),  
            nn.ReLU(),  
            nn.Linear(512, 10)  
        )
```

Create Neural Networks Model

```
    def forward(self, x):  
        x = self.flatten(x)  
        logits = self.linear_relu_stack(x)  
        return logits
```

```
model = NeuralNetwork().to(device)  
print(model)
```

Optimizing the Model Parameters

```
loss_fn = nn.CrossEntropyLoss()  
optimizer = torch.optim.SGD(model.parameters(), lr=1e-3)
```

Define Training Loop

```
def train(dataloader, model, loss_fn, optimizer):
    size = len(dataloader.dataset)
    model.train()
    for batch, (X, y) in enumerate(dataloader):
        X, y = X.to(device), y.to(device)

        # Compute prediction error
        pred = model(X)
        loss = loss_fn(pred, y)

        # Backpropagation
        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

    if batch % 100 == 0:
        loss, current = loss.item(), batch * len(X)
        print(f"loss: {loss:>7f}  [{current:>5d}/{size:>5d}"])
```

Define Test Dataset

```
def test(dataloader, model, loss_fn):
    size = len(dataloader.dataset)
    num_batches = len(dataloader)
    model.eval()
    test_loss, correct = 0, 0
    with torch.no_grad():
        for X, y in dataloader:
            X, y = X.to(device), y.to(device)
            pred = model(X)
            test_loss += loss_fn(pred, y).item()
            correct += (pred.argmax(1) == y).type(torch.float).sum().item()
    test_loss /= num_batches
    correct /= size
    print(f"Test Error: \n Accuracy: {(100*correct):>0.1f}%, Avg loss: {test_loss:>8f} \n")
```


Training Loop

```
epochs = 5
for t in range(epochs):
    print(f"Epoch {t+1}\n-----")
    train(train_dataloader, model, loss_fn, optimizer)
    test(test_dataloader, model, loss_fn)
print("Done!")
```

Saving Models

```
torch.save(model.state_dict(), "model.pth")  
print("Saved PyTorch Model State to model.pth")
```

Loading Models

```
model = NeuralNetwork()  
model.load_state_dict(torch.load("model.pth"))
```

Make Predictions

```
classes = [  
    "T-shirt/top",  
    "Trouser",  
    "Pullover",  
    "Dress",  
    "Coat",  
    "Sandal",  
    "Shirt",  
    "Sneaker",  
    "Bag",  
    "Ankle boot",  
]  
  
model.eval()  
x, y = test_data[0][0], test_data[0][1]  
with torch.no_grad():  
    pred = model(x)  
    predicted, actual = classes[pred[0].argmax(0)], classes[y]  
    print(f'Predicted: "{predicted}", Actual: "{actual}"')
```